

Package ‘convo’

June 18, 2018

Type Package

Title Fast Computation of Running Sample Statistics

Version 0.1.0

Description Provides methods for fast computation of running (aka rolling) sample statistics. These include running sample mean of a single numeric sequence, running sample variance of a single numeric sequence, running sample covariance of two numeric sequences. Implementation of the methods uses convolution via Fast Fourier Transform. From Convolution Theorem, the complexity of the convolution can be reduced from $O(n^2)$ to $O(n \log n)$ with fast Fourier transform compared to conventional computation.

License GPL-3

Imports stats

Encoding UTF-8

LazyData true

RoxygenNote 6.0.1

Suggests knitr, rmarkdown, testthat, covr

VignetteBuilder knitr

Author Marta Karas [aut, cre],
Jacek Urbanek [aut]

Maintainer Marta Karas <mkaras2@jhu.edu>

R topics documented:

convJU	2
DigFilter	3
RunningCorr	3
RunningCov	4
RunningL2Norm	5
RunningMean	6
RunningSd	7
RunningVar	8

Index	9
--------------	----------

convJU

*Fast Sequences Convolution***Description**

Computes convolutions of two sequences via Fast Fourier Transform

Usage

```
convJU(x, y)
```

Arguments

x	numeric sequence
y	numeric sequence, of equal or shorter length than x sequence

Details

Use the Fast Fourier Transform to compute convolutions of two sequences. If sequences are of different length, the shorter one get a suffix of 0's. Following convention of `stats::convolve` function, if `r <- convolve(x, y, type = "open")` and `n <- length(x)`, `m <- length(y)`, then

$$r[k] = \sum_i x[k - m + i] \cdot y[i]$$

where the sum is over all valid indices i .

FFT formulation is useful for implementing an efficient numerical convolution: the standard convolution algorithm has quadratic computational complexity. From convolution theorem, the complexity of the convolution can be reduced from $O(n^2)$ to $O(n \log n)$ with fast Fourier transform.

Value

numeric sequence

Examples

```
## Not run:
x <- sin(seq(0, 1, length.out = 1000) * 2 * pi * 6)
y <- rep(1, 100)
convJU(x, y)

## End(Not run)
```

DigFilter

*Bandpass Digital Filter***Description**

Passes frequencies within a certain range and rejects (attenuates) frequencies outside that range.

Usage

```
DigFilter(x, fs, LD, LU)
```

Arguments

x	numeric sequence
fs	sampling frequency of numeric sequence x
LD	lower bandpass frequency
LU	higher bandpass frequency

Examples

```
## generate components of signal x
fs <- 1000
t.seq <- seq(0, 1, length.out = 1000)
x1 <- sin(2 * pi * t.seq * 50)
x2 <- sin(2 * pi * t.seq * 150)
x3 <- sin(2 * pi * t.seq * 250)
x0 <- rnorm(length(t.seq), sd = 0.1)
plot(x0, ylim = range(c(x0, x1, x2, x3)), main = "x components")
lines(x1, col = "blue")
lines(x2, col = "red")
lines(x3, col = "green")
## generate signal x
x <- x1 + x2 + x3 + x0
plot(x, main = "x", type = "l")
## use filter with passband frequencies [100 Hz, 200 Hz]
x.filtered <- DigFilter(x, 1000, 100, 200)
plot(x.filtered, type = "l", main = "x filtered with [100 Hz, 200 Hz] frequency bandpass")
plot(x2, type = "l", main = "compare: x2 component of x")
```

RunningCorr

*Fast Running Correlation Computation***Description**

Computes running correlation between two sequences in a fixed width window, whose length corresponds to the length of the shorter sequence. Uses convolution via Fast Fourier Transform.

Usage

```
RunningCorr(x, y, circular = FALSE)
```

Arguments

x	numeric sequence
y	numeric sequence, of equal or shorter length than x sequence
circular	logical; whether running correlation is computed assuming circular nature of x sequence (see Details)

Details

Computes running correlation between two sequences in a fixed width window. The length of a window is equal to the shorter of the two sequences (y), and window "runs" over the length of longer sequence (x).

Parameter circular determines whether x sequence is assumed to have a circular nature. Assume l_x is the length of longer sequence x, l_y is the length of shorter sequence y. * If circular equals TRUE, then the output sequence has length equal to the length of x sequence. First element of the output sequence corresponds to correlation between $x[1:l_y]$ window of x sequence and y sequence. Last element of the output sequence corresponds to correlation between $x[(l_x):(l_y - 1)]$ window of x sequence and y sequence. * If circular equals FALSE, then the output sequence has length equal to $l_x - l_y + 1$. First element of the output sequence corresponds to correlation between $x[1:l_y]$ window of x sequence and y sequence. Last element of the output sequence corresponds to correlation between $x[(l_x - l_y + 1):l_x]$ window of x sequence and y sequence.

Value

numeric sequence

Examples

```
x <- sin(seq(0, 1, length.out = 1000) * 2 * pi * 6)
y <- x[1:100]
out1 <- RunningCorr(x, y, circular = TRUE)
out2 <- RunningCorr(x, y, circular = FALSE)
plot(out1, type = "l"); points(out2, col = "red")
```

RunningCov

Fast Running Covariance Computation

Description

Computes running covariance between two sequences in a fixed width window, whose length corresponds to the length of the shorter sequence. Uses convolution via Fast Fourier Transform.

Usage

```
RunningCov(x, y, circular = FALSE)
```

Arguments

x	numeric sequence
y	numeric sequence, of equal or shorter length than x sequence
circular	logical; whether running variance is computed assuming circular nature of x sequence (see Details)

Details

Computes running covariance between two sequences in a fixed width window. The length of a window is equal to the shorter of the two sequences (y), and window "runs" over the length of longer sequence (x).

Parameter `circular` determines whether x sequence is assumed to have a circular nature. Assume l_x is the length of longer sequence x , l_y is the length of shorter sequence y . * If `circular` equals `TRUE`, then the output sequence has length equal to the length of x sequence. First element of the output sequence corresponds to covariance between $x[1:l_y]$ window of x sequence and y sequence. Last element of the output sequence corresponds to covariance between $x[(l_x):(l_y - 1)]$ window of x sequence and y sequence. * If `circular` equals `FALSE`, then the output sequence has length equal to $l_x - l_y + 1$. First element of the output sequence corresponds to covariance between $x[1:l_y]$ window of x sequence and y sequence. Last element of the output sequence corresponds to covariance between $x[(l_x - l_y + 1):l_x]$ window of x sequence and y sequence.

Value

numeric sequence

Examples

```
x <- sin(seq(0, 1, length.out = 1000) * 2 * pi * 6)
y <- x[1:100]
out1 <- RunningCov(x, y, circular = TRUE)
out2 <- RunningCov(x, y, circular = FALSE)
plot(out1, type = "l"); points(out2, col = "red")
```

RunningL2Norm

Fast Running L2 Norm Computation

Description

Computes running L2 norm between two sequences in a fixed width window, whose length corresponds to the length of the shorter sequence. Uses convolution via Fast Fourier Transform.

Usage

```
RunningL2Norm(x, y, circular = FALSE)
```

Arguments

<code>x</code>	numeric sequence
<code>y</code>	numeric sequence, of equal or shorter length than x sequence
<code>circular</code>	logical; whether running L2 norm is computed assuming circular nature of x sequence (see Details)

Details

Computes running L2 norm between two sequences in a fixed width window. The length of a window is equal to the shorter of the two sequences (y), and window "runs" over the length of longer sequence (x).

Parameter `circular` determines whether x sequence is assumed to have a circular nature. Assume l_x is the length of longer sequence x , l_y is the length of shorter sequence y . * If `circular` equals `TRUE`, then the output sequence has length equal to the length of x sequence. First element of the output sequence corresponds to L2 norm between $x[1:l_y]$ window of x sequence and y sequence. Last element of the output sequence corresponds to L2 norm between $x[(l_x - l_y + 1):l_x]$ window of x sequence and y sequence. * If `circular` equals `FALSE`, then the output sequence has length equal to $l_x - l_y + 1$. First element of the output sequence corresponds to L2 norm between $x[1:l_y]$ window of x sequence and y sequence. Last element of the output sequence corresponds to L2 norm between $x[(l_x - l_y + 1):l_x]$ window of x sequence and y sequence.

Value

numeric sequence

Examples

```
x <- sin(seq(0, 1, length.out = 1000) * 2 * pi * 6)
y1 <- x[1:100] + rnorm(100)
y2 <- rnorm(100)
out1 <- RunningL2Norm(x, y1)
out2 <- RunningL2Norm(x, y2)
plot(out1, type = "l"); points(out2, col = "blue")
```

RunningMean

Fast Running Mean Computation

Description

Computes running sample mean of a sequence in a fixed width window. Uses convolution via Fast Fourier Transform.

Usage

```
RunningMean(x, W, circular = FALSE)
```

Arguments

<code>x</code>	numeric sequence
<code>W</code>	numeric; width of x sequence window
<code>circular</code>	logical; whether running sample mean is computed assuming circular nature of x sequence (see Details)

Details

Parameter `circular` determines whether `x` sequence is assumed to have a circular nature. Assume l_x is the length of sequence `x`, `W` is a fixed length of codex sequence window. * If `circular` equals `TRUE`, then the output sequence has length equal to the length of `x` sequence. First element of the output sequence corresponds to sample mean of `x[1:W]` window of `x` sequence. Last element of the output sequence corresponds to sample mean of `x[(1_x):(W - 1)]` window of `x` sequence. * If `circular` equals `FALSE`, then the output sequence has length equal to $l_x - W + 1$. First element of the output sequence corresponds to sample mean of `x[1:W]` window of `x` sequence. Last element of the output sequence corresponds to sample mean of `x[(1_x - W + 1):1_x]` window of `x` sequence.

Value

numeric sequence

Examples

```
x <- rnorm(1000)
RunningMean(x, 100)
length(RunningMean(x, 100, circular = FALSE))
length(RunningMean(x, 100, circular = TRUE))
```

RunningSd

Fast Running Standard Deviation Computation

Description

Computes running sample standard deviation of a sequence in a fixed width window. Uses convolution via Fast Fourier Transform.

Usage

```
RunningSd(x, W, circular = FALSE)
```

Arguments

<code>x</code>	numeric sequence
<code>W</code>	numeric; width of <code>x</code> sequence window
<code>circular</code>	logical; whether running sample standard deviation is computed assuming circular nature of <code>x</code> sequence (see Details)

Details

Parameter `circular` determines whether `x` sequence is assumed to have a circular nature. Assume l_x is the length of sequence `x`, `W` is a fixed length of codex sequence window. * If `circular` equals `TRUE`, then the output sequence has length equal to the length of `x` sequence. First element of the output sequence corresponds to sample standard deviation of `x[1:W]` window of `x` sequence. Last element of the output sequence corresponds to sample standard deviation of `x[(1_x):(W - 1)]` window of `x` sequence. * If `circular` equals `FALSE`, then the output sequence has length equal to $l_x - W + 1$. First element of the output sequence corresponds to sample standard deviation of `x[1:W]` window of `x` sequence. Last element of the output sequence corresponds to sample standard deviation of `x[(1_x - W + 1):1_x]` window of `x` sequence.

Value

numeric sequence

Examples

```
x <- rnorm(1000)
RunningSd(x, 100)
length(RunningSd(x, 100, circular = FALSE))
length(RunningSd(x, 100, circular = TRUE))
```

RunningVar

Fast Running Variance Computation

Description

Computes running sample variance of a sequence in a fixed width window. Uses convolution via Fast Fourier Transform.

Usage

```
RunningVar(x, W, circular = FALSE)
```

Arguments

x	numeric sequence
W	numeric; width of x sequence window
circular	logical; whether running sample variance is computed assuming circular nature of x sequence (see Details)

Details

Parameter `circular` determines whether x sequence is assumed to have a circular nature. Assume l_x is the length of sequence x, W is a fixed length of codex sequence window. * If `circular` equals TRUE, then the output sequence has length equal to the length of x sequence. First element of the output sequence corresponds to sample variance of $x[1:W]$ window of x sequence. Last element of the output sequence corresponds to sample variance of $x[(l_x):(W - 1)]$ window of x sequence. * If `circular` equals FALSE, then the output sequence has length equal to $l_x - W + 1$. First element of the output sequence corresponds to sample variance of $x[1:W]$ window of x sequence. Last element of the output sequence corresponds to sample variance of $x[(l_x - W + 1):l_x]$ window of x sequence.

Value

numeric sequence

Examples

```
x <- rnorm(1000)
RunningVar(x, 100)
length(RunningVar(x, 100, circular = FALSE))
length(RunningVar(x, 100, circular = TRUE))
```


Index

convJU, [2](#)

DigFilter, [3](#)

RunningCorr, [3](#)

RunningCov, [4](#)

RunningL2Norm, [5](#)

RunningMean, [6](#)

RunningSd, [7](#)

RunningVar, [8](#)